

Assignment4

March 22, 2025

```
[71]: import pandas as pd
import numpy as np
```

```
[72]: df = pd.read_csv('sales.csv' , encoding = 'latin-1')
```

```
[73]: df.head()
```

```
[73]:  ORDERNUMBER  QUANTITYORDERED  PRICEEACH  ORDERLINENUMBER  SALES  \
0          10107                30        95.70                2  2871.00
1          10121                34        81.35                5  2765.90
2          10134                41        94.74                2  3884.34
3          10145                45        83.26                6  3746.70
4          10159                49       100.00               14  5205.27
```

```
  ORDERDATE  STATUS  QTR_ID  MONTH_ID  YEAR_ID  ...  \
0  2/24/2003 0:00  Shipped        1         2    2003  ...
1  5/7/2003 0:00  Shipped        2         5    2003  ...
2  7/1/2003 0:00  Shipped        3         7    2003  ...
3  8/25/2003 0:00  Shipped        3         8    2003  ...
4 10/10/2003 0:00  Shipped        4        10    2003  ...
```

```
  ADDRESSLINE1  ADDRESSLINE2  CITY STATE  \
0  897 Long Airport Avenue    NaN    NYC   NY
1      59 rue de l'Abbaye    NaN   Reims  NaN
2  27 rue du Colonel Pierre Avia    NaN   Paris  NaN
3      78934 Hillside Dr.    NaN  Pasadena   CA
4      7734 Strong St.    NaN San Francisco   CA
```

```
  POSTALCODE  COUNTRY  TERRITORY  CONTACTLASTNAME  CONTACTFIRSTNAME  DEALSIZE
0       10022     USA      NaN           Yu           Kwai      Small
1       51100  France     EMEA      Henriot           Paul      Small
2       75508  France     EMEA    Da Cunha      Daniel      Medium
3       90003     USA      NaN      Young           Julie      Medium
4         NaN     USA      NaN      Brown           Julie      Medium
```

```
[5 rows x 25 columns]
```

```
[74]: df.describe(include='all')
```

```
[74]:
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	\
count	2823.000000	2823.000000	2823.000000	2823.000000	
unique	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	
mean	10258.725115	35.092809	83.658544	6.466171	
std	92.085478	9.741443	20.174277	4.225841	
min	10100.000000	6.000000	26.880000	1.000000	
25%	10180.000000	27.000000	68.860000	3.000000	
50%	10262.000000	35.000000	95.700000	6.000000	
75%	10333.500000	43.000000	100.000000	9.000000	
max	10425.000000	97.000000	100.000000	18.000000	

	SALES	ORDERDATE	STATUS	QTR_ID	MONTH_ID	\
count	2823.000000	2823	2823	2823.000000	2823.000000	
unique	NaN	252	6	NaN	NaN	
top	NaN	11/14/2003 0:00	Shipped	NaN	NaN	
freq	NaN	38	2617	NaN	NaN	
mean	3553.889072	NaN	NaN	2.717676	7.092455	
std	1841.865106	NaN	NaN	1.203878	3.656633	
min	482.130000	NaN	NaN	1.000000	1.000000	
25%	2203.430000	NaN	NaN	2.000000	4.000000	
50%	3184.800000	NaN	NaN	3.000000	8.000000	
75%	4508.000000	NaN	NaN	4.000000	11.000000	
max	14082.800000	NaN	NaN	4.000000	12.000000	

	YEAR_ID	...	ADDRESSLINE1	ADDRESSLINE2	CITY	STATE	\
count	2823.000000	...	2823	302	2823	1337	
unique	NaN	...	92	9	73	16	
top	NaN	...	C/ Moralzarzal, 86	Level 3	Madrid	CA	
freq	NaN	...	259	55	304	416	
mean	2003.81509	...	NaN	NaN	NaN	NaN	
std	0.69967	...	NaN	NaN	NaN	NaN	
min	2003.000000	...	NaN	NaN	NaN	NaN	
25%	2003.000000	...	NaN	NaN	NaN	NaN	
50%	2004.000000	...	NaN	NaN	NaN	NaN	
75%	2004.000000	...	NaN	NaN	NaN	NaN	
max	2005.000000	...	NaN	NaN	NaN	NaN	

	POSTALCODE	COUNTRY	TERRITORY	CONTACTLASTNAME	CONTACTFIRSTNAME	DEALSIZE
count	2747	2823	1749	2823	2823	2823
unique	73	19	3	77	72	3
top	28034	USA	EMEA	Freyre	Diego	Medium
freq	259	1004	1407	259	259	1384
mean	NaN	NaN	NaN	NaN	NaN	NaN

std	NaN	NaN	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN	NaN	NaN

[11 rows x 25 columns]

```
[75]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ORDERNUMBER           2823 non-null   int64
1   QUANTITYORDERED       2823 non-null   int64
2   PRICEEACH             2823 non-null   float64
3   ORDERLINENUMBER       2823 non-null   int64
4   SALES                 2823 non-null   float64
5   ORDERDATE             2823 non-null   object
6   STATUS                2823 non-null   object
7   QTR_ID               2823 non-null   int64
8   MONTH_ID              2823 non-null   int64
9   YEAR_ID               2823 non-null   int64
10  PRODUCTLINE           2823 non-null   object
11  MSRP                  2823 non-null   int64
12  PRODUCTCODE           2823 non-null   object
13  CUSTOMERNAME          2823 non-null   object
14  PHONE                 2823 non-null   object
15  ADDRESSLINE1          2823 non-null   object
16  ADDRESSLINE2          302 non-null    object
17  CITY                  2823 non-null   object
18  STATE                 1337 non-null   object
19  POSTALCODE            2747 non-null   object
20  COUNTRY               2823 non-null   object
21  TERRITORY             1749 non-null   object
22  CONTACTLASTNAME       2823 non-null   object
23  CONTACTFIRSTNAME      2823 non-null   object
24  DEALSIZE              2823 non-null   object
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

```
[76]: df.shape
```

```
[76]: (2823, 25)
```


```
[77]: df.isnull().sum()
```

```
[77]: ORDERNUMBER          0
      QUANTITYORDERED    0
      PRICEEACH          0
      ORDERLINENUMBER    0
      SALES              0
      ORDERDATE          0
      STATUS            0
      QTR_ID            0
      MONTH_ID          0
      YEAR_ID           0
      PRODUCTLINE       0
      MSRP              0
      PRODUCTCODE       0
      CUSTOMERNAME      0
      PHONE             0
      ADDRESSLINE1       0
      ADDRESSLINE2      2521
      CITY              0
      STATE             1486
      POSTALCODE         76
      COUNTRY           0
      TERRITORY         1074
      CONTACTLASTNAME    0
      CONTACTFIRSTNAME   0
      DEALSIZE           0
      dtype: int64
```


```
[78]: df = df.drop('STATE', axis=1)
      df = df.drop('POSTALCODE', axis=1)
      df = df.drop('TERRITORY', axis=1)
      df = df.drop('ADDRESSLINE2', axis=1)
```

```
[79]: import matplotlib.pyplot as plt
      import seaborn as sns
```

```
[80]: sns.histplot(df["SALES"], bins=30, kde=True, color="blue")
      plt.title("Histogram of SALES")
      plt.xlabel("Sales Amount")
      plt.ylabel("Frequency")
      plt.show()
```



```
[81]: sns.histplot(x=df["MONTH_ID"], y=df["SALES"], bins=12, kde=True, cmap="Blues")
plt.title("Histogram of SALES vs. MONTH_ID")
plt.xlabel("Month ID")
plt.ylabel("Sales Amount")
plt.xticks(range(1, 13))
plt.show()
```

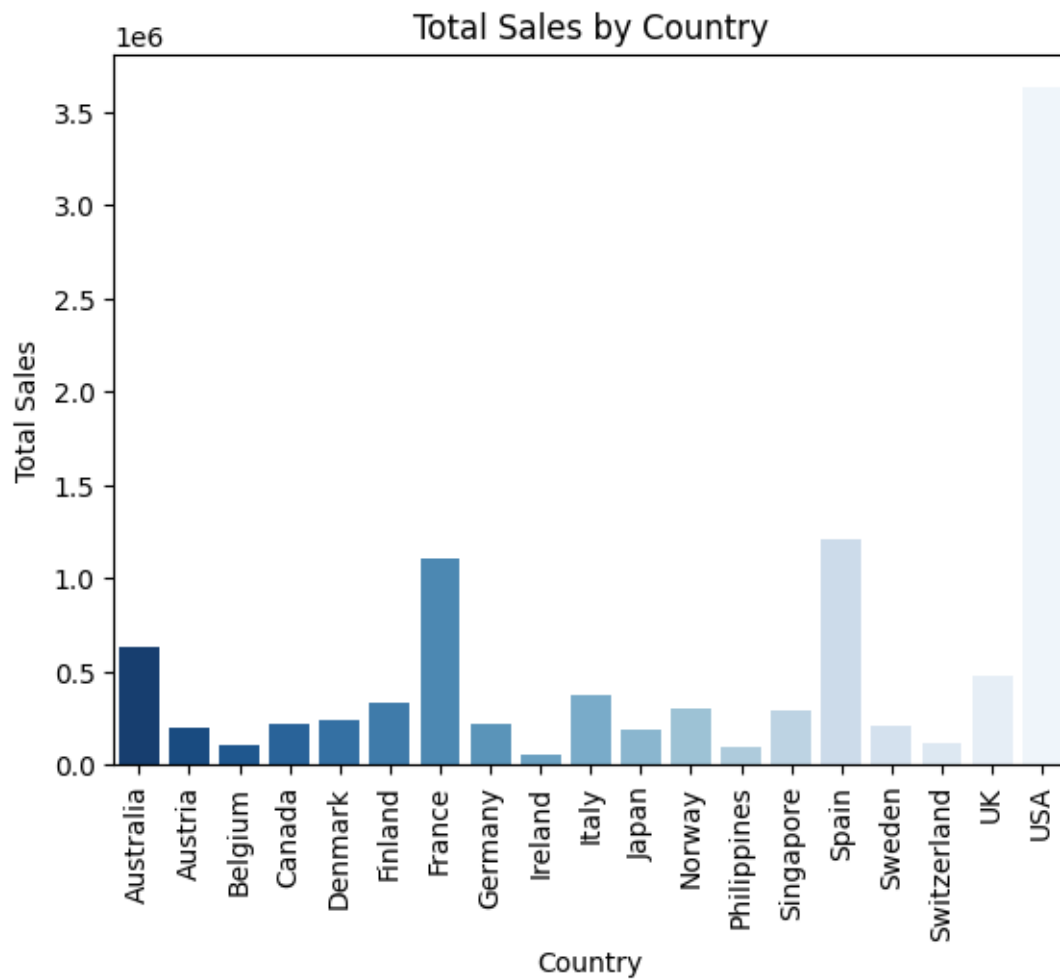


```
[82]: sns.barplot(x=df.groupby("COUNTRY")["SALES"].sum().index,
                  y=df.groupby("COUNTRY")["SALES"].sum().values,
                  palette="Blues_r")
plt.xticks(rotation=90)
plt.title("Total Sales by Country")
plt.xlabel("Country")
plt.ylabel("Total Sales")
plt.show()
```


<ipython-input-82-6da0283484d0>:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.


```
sns.barplot(x=df.groupby("COUNTRY")["SALES"].sum().index,
```




```
[83]: sns.histplot(df["ORDERNUMBER"], bins=30, kde=True, color="teal")
plt.title("Histogram of ORDERNUMBER")
plt.xlabel("Order Number")
plt.ylabel("Frequency")
plt.show()
```




```
[84]: sns.histplot(df["SALES"], bins=30, kde=True, color="blue")
plt.title("Histogram of SALES")
plt.xlabel("Sales Amount")
plt.ylabel("Frequency")
plt.show()
```


```
[85]: sns.histplot(df["MSRP"], bins=30, kde=True, color="green")  
plt.title("Histogram of MSRP")  
plt.xlabel("MSRP")  
plt.ylabel("Frequency")  
plt.show()
```



```
[86]: sns.kdeplot(df["SALES"], fill=True, color="blue")
plt.title("KDE Plot of SALES")
plt.xlabel("Sales Amount")
plt.ylabel("Density")
plt.show()
```




```
[87]: sns.kdeplot(df["QUANTITYORDERED"], fill=True, color="red")  
plt.title("KDE Plot of QUANTITYORDERED")  
plt.xlabel("Quantity Ordered")  
plt.ylabel("Density")  
plt.show()
```




```
[88]: plt.figure(figsize=(6, 6))
df["STATUS"].value_counts().plot.pie(autopct="%1.1f%%", colors=["skyblue", "lightcoral", "lightgreen", "orange", "purple"])
plt.title("Order Status Distribution")
plt.ylabel("")
plt.show()
```

Order Status Distribution



```
[89]: plt.figure(figsize=(6, 6))
df["DEALSIZE"].value_counts().plot.pie(autopct="%1.1f%%", colors=["skyblue", "lightcoral", "lightgreen"], startangle=90)
plt.title("Deal Size Distribution")
plt.ylabel("")
plt.show()
```



- Generate a small dataset from a non-normal distribution (e.g., uniform distribution).

```
[90]: uniform_data = np.random.uniform(low=10, high=50, size=100)
```


```
[91]: df_uniform = pd.DataFrame(uniform_data, columns=["Uniform_Distribution"])
```

```
[92]: df_uniform.head()
```

```
[92]:    Uniform_Distribution
0          36.052119
1          44.958219
2          36.071830
3          44.038639
4          28.605493
```

```
[93]: plt.figure(figsize=(7, 5))
```

```
sns.histplot(df_uniform["Uniform_Distribution"], bins=10, kde=True,
             color="blue")
plt.title("Histogram of Uniform Distribution")
plt.xlabel("Value")
plt.ylabel("Frequency")
plt.show()
```



- Take multiple random samples from the dataset and calculate their means.

```
[94]: num_samples = 30
      sample_size = 10
```

```
[95]: sample_means = []


for _ in range(num_samples):
    sample = df_uniform["Uniform_Distribution"].sample(sample_size,
                                                         replace=True)
    sample_means.append(sample.mean())
```

```
[96]: df_sample_means = pd.DataFrame(sample_means, columns=["Sample_Means"])
```

```
[97]: df_sample_means.head()
```

```
[97]: Sample_Means
0      31.461911
1      25.024023
2      26.445996
3      24.220957
4      31.559834
```

```
[98]: plt.figure(figsize=(7, 5))
sns.histplot(df_sample_means["Sample_Means"], bins=10, kde=True, color="green")
plt.title(f"Histogram of Sample Means (Sample Size = {sample_size})")
plt.xlabel("Sample Mean")
plt.ylabel("Frequency")
plt.show()
```



Hypothesis Testing :

```
[99]: from scipy.stats import ttest_ind
```



```

sales_small = df[df["DEALSIZE"] == "Small"]["SALES"]
sales_large = df[df["DEALSIZE"] == "Large"]["SALES"]

t_stat, p_value = ttest_ind(sales_small, sales_large, equal_var=False)

print(f"T-statistic: {t_stat:.4f}")
print(f"P-value: {p_value:.4f}")

alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant difference in_
↪SALES between Small and Large deal sizes.")
else:
    print("Fail to reject the null hypothesis: No significant difference in_
↪SALES between Small and Large deal sizes.")

```

T-statistic: -59.5907

P-value: 0.0000

Reject the null hypothesis: There is a significant difference in SALES between Small and Large deal sizes.

```

[100]: contingency_table = pd.crosstab(df["STATUS"], df["DEALSIZE"])

chi2_stat, p_value, dof, expected = stats.chi2_contingency(contingency_table)

print(f"Chi-Square Statistic: {chi2_stat:.4f}")
print(f"Degrees of Freedom: {dof}")
print(f"P-value: {p_value:.4f}")

alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: STATUS and DEALSIZE are dependent.")
else:
    print("Fail to reject the null hypothesis: No significant relationship_
↪between STATUS and DEALSIZE.")

```

Chi-Square Statistic: 34.3444

Degrees of Freedom: 10

P-value: 0.0002

Reject the null hypothesis: STATUS and DEALSIZE are dependent.

```

[101]: len(sales_data)

```

[101]: 2823

```

[102]: mean_sales = np.mean(sales_data)
std_error = stats.sem(sales_data)

```

```
[103]: mean_sales
```

```
[103]: np.float64(3553.889071909316)
```

```
[104]: std_error
```

```
[104]: np.float64(34.66589211924902)
```

```
[105]: confidence = 0.95
margin_of_error = stats.t.ppf((1 + confidence) / 2, len(sales_data) - 1) *
    ↪std_error

lower_bound = mean_sales - margin_of_error
upper_bound = mean_sales + margin_of_error

print(f"95% Confidence Interval for Mean SALES: ({lower_bound:.2f},
    ↪{upper_bound:.2f})")

print("Interpretation: We are 95% confident that the true mean SALES value lies
    ↪within this range.")
```

95% Confidence Interval for Mean SALES: (3485.92, 3621.86)

Interpretation: We are 95% confident that the true mean SALES value lies within this range.

```
[106]: import scipy.stats as stats

usa_sales = df[df["COUNTRY"] == "USA"]["SALES"].dropna()
france_sales = df[df["COUNTRY"] == "France"]["SALES"].dropna()

t_stat, p_value = stats.ttest_ind(usa_sales, france_sales, equal_var=False) #
    ↪Welch's t-test

print(f"T-Statistic: {t_stat:.4f}")
print(f"P-Value: {p_value:.4f}")

alpha = 0.05
if p_value < alpha:
    print("Result: Reject the Null Hypothesis (Significant difference in SALES
    ↪between USA and France)")
else:
    print("Result: Fail to Reject the Null Hypothesis (No significant
    ↪difference in SALES between USA and France)")
```

T-Statistic: 0.6063

P-Value: 0.5446

Result: Fail to Reject the Null Hypothesis (No significant difference in SALES

between USA and France)

```
[107]: from scipy.stats import skew, kurtosis

sales_skewness = skew(df["SALES"], nan_policy='omit')
print(f"Skewness of SALES: {sales_skewness:.4f}")

sales_kurtosis = kurtosis(df["SALES"], nan_policy='omit')
print(f"Kurtosis of SALES: {sales_kurtosis:.4f}")
```

Skewness of SALES: 1.1605

Kurtosis of SALES: 1.7874

```
[108]: import pandas as pd
from scipy.stats import skew, kurtosis

numerical_cols = df.select_dtypes(include=['number'])

skewness_values = numerical_cols.apply(lambda x: skew(x, nan_policy='omit'))
kurtosis_values = numerical_cols.apply(lambda x: kurtosis(x, nan_policy='omit'))

stats_df = pd.DataFrame({'Skewness': skewness_values, 'Kurtosis':
    ↳kurtosis_values})
print(stats_df)
```

	Skewness	Kurtosis
ORDERNUMBER	0.013816	-1.173357
QUANTITYORDERED	0.362393	0.412883
PRICEEACH	-0.946146	-0.376279
ORDERLINENUMBER	0.590427	-0.562285
SALES	1.160459	1.787378
QTR_ID	-0.255815	-1.498237
MONTH_ID	-0.272757	-1.382951
YEAR_ID	0.271307	-0.951001
MSRP	0.579867	-0.133706

```
[109]: highly_skewed_cols = stats_df[stats_df['Skewness'].abs() > 1].index

for i, col in enumerate(highly_skewed_cols):
    plt.subplot(len(highly_skewed_cols), 2, 2*i+1)
    sns.histplot(df[col], bins=30, kde=True, color="skyblue")
    plt.title(f"Histogram & KDE of {col}")

    plt.subplot(len(highly_skewed_cols), 2, 2*i+2)
    sns.boxplot(x=df[col], color="lightcoral")
    plt.title(f"Boxplot of {col}")

plt.tight_layout()
```

```
plt.show()
```

