

Assignment6

March 21, 2025

EDA

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: df = pd.read_csv('churn.csv')
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   phoneno     5000 non-null   int64
 1   age         4994 non-null   float64
 2   gender      5000 non-null   object
 3   zipcode     5000 non-null   int64
 4   calls       5000 non-null   int64
 5   sms         5000 non-null   int64
 6   mms         5000 non-null   int64
 7   charges     5000 non-null   int64
 8   coverage    5000 non-null   int64
 9   complaint   5000 non-null   int64
10   sim         5000 non-null   object
11   phone       5000 non-null   object
12   prepost     5000 non-null   object
13   churn       5000 non-null   object
dtypes: float64(1), int64(8), object(5)
memory usage: 547.0+ KB
```

```
[ ]: df.head()
```

```
[ ]:   phoneno  age  gender  zipcode  calls  sms  mms  charges  coverage  \
0     5974  1.0   Male    91107    160   25    1      490         0
1     4535  1.0   Male    90089    150   45   19      340         0
```

2	4016	1.0	Male	94720	100	39	15	110	0
3	8523	2.0	Male	94112	270	35	9	1000	0
4	5052	2.0	Female	91330	100	35	8	450	0

	complaint	sim	phone	prepost	churn
0	4	Dual Sim	Andoid	Prepaid	No Churn
1	3	Dual Sim	Andoid	Prepaid	No Churn
2	1	Single Sim	Andoid	Prepaid	No Churn
3	1	Single Sim	Andoid	Prepaid	No Churn
4	4	Single Sim	Andoid	Prepaid	No Churn

```
[ ]: df.columns
```

```
[ ]: Index(['phoneno', 'age', 'gender', 'zipcode', 'calls', 'sms', 'mms', 'charges',
           'coverage', 'complaint', 'sim', 'phone', 'prepost', 'churn'],
          dtype='object')
```

```
[ ]: df.describe(include='all')
```

```
[ ]:
```

	phoneno	age	gender	zipcode	calls \
count	5000.000000	4994.000000	5000	5000.000000	5000.000000
unique	NaN	NaN	2	NaN	NaN
top	NaN	NaN	Male	NaN	NaN
freq	NaN	NaN	3530	NaN	NaN
mean	5497.188000	1.881057	NaN	93152.503000	193.793800
std	2603.474018	0.839796	NaN	2121.852197	174.765898
min	1000.000000	1.000000	NaN	9307.000000	0.000000
25%	3266.500000	1.000000	NaN	91911.000000	70.000000
50%	5457.500000	2.000000	NaN	93437.000000	150.000000
75%	7779.250000	3.000000	NaN	94608.000000	250.000000
max	9997.000000	3.000000	NaN	96651.000000	1000.000000

	sms	mms	charges	coverage	complaint \
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000
unique	NaN	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN	NaN
mean	45.338400	20.104600	737.742000	0.719000	2.388600
std	11.463166	11.467954	460.337293	1.233184	1.154061
min	23.000000	-3.000000	80.000000	0.000000	0.000000
25%	35.000000	10.000000	390.000000	0.000000	1.000000
50%	45.000000	20.000000	640.000000	0.000000	2.000000
75%	55.000000	30.000000	980.000000	2.000000	3.000000
max	67.000000	43.000000	2240.000000	7.000000	4.000000

	sim	phone	prepost	churn
count	5000	5000	5000	5000

unique	2	2	2	2
top	Single Sim	Andoid	Postpaid	No Churn
freq	4478	4698	2984	4520
mean	NaN	NaN	NaN	NaN
std	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN

```
[ ]: df.shape
```

```
[ ]: (5000, 14)
```

```
[ ]: df.tail()
```

```
[ ]:
      phoneno  age  gender  zipcode  calls  sms  mms  charges  coverage \
4995     4704  3.0   Male    92697    190   29   3     400         0
4996     3149  1.0   Male    92037     40   30   4     150         1
4997     7402  3.0   Male    93023     30   63  39     240         0
4998     5742  2.0   Male    90034     50   65  40     490         0
4999     3689  1.0  Female    92612     80   28   4     830         0
```

	complaint	sim	phone	prepost	churn
4995	1	Single Sim	Andoid	Postpaid	No Churn
4996	4	Single Sim	Andoid	Postpaid	No Churn
4997	2	Single Sim	Andoid	Prepaid	No Churn
4998	3	Single Sim	Andoid	Postpaid	No Churn
4999	3	Single Sim	Andoid	Postpaid	No Churn

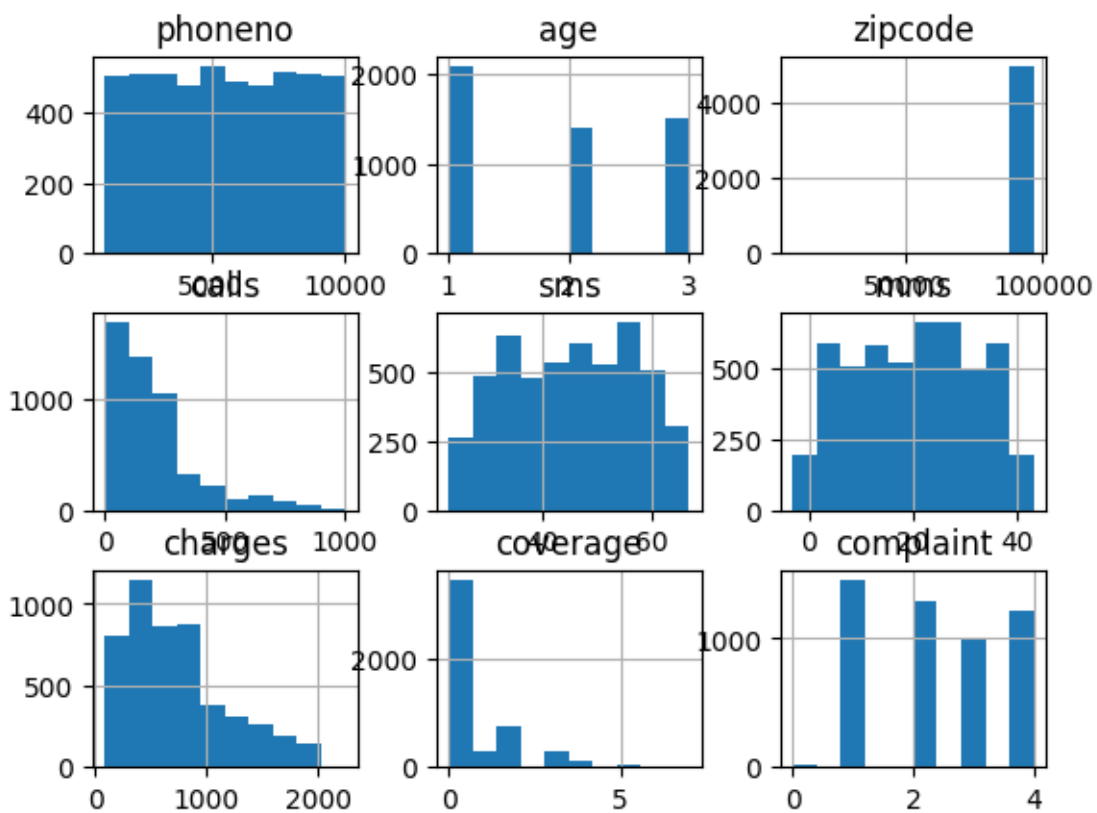
```
[ ]: df.isnull().sum()
```

```
[ ]: phoneno      0
      age         6
      gender      0
      zipcode     0
      calls       0
      sms         0
      mms         0
      charges     0
      coverage    0
      complaint   0
      sim         0
      phone       0
      prepost     0
      churn       0
```

dtype: int64


```
[ ]: df.hist()
```

```
[ ]: array([[<Axes: title={'center': 'phoneno'}>,  
          <Axes: title={'center': 'age'}>,  
          <Axes: title={'center': 'zipcode'}>],  
          [<Axes: title={'center': 'calls'}>,  
          <Axes: title={'center': 'sms'}>, <Axes: title={'center': 'mms'}>],  
          [<Axes: title={'center': 'charges'}>,  
          <Axes: title={'center': 'coverage'}>,  
          <Axes: title={'center': 'complaint'}>]], dtype=object)
```




```
[ ]: plt.hist(df['complaint'])
```

```
[ ]: (array([ 18.,   0., 1464.,   0.,   0., 1292.,   0., 1009.,   0.,  
          1217.]),  
      array([0. , 0.4, 0.8, 1.2, 1.6, 2. , 2.4, 2.8, 3.2, 3.6, 4. ]),  
      <BarContainer object of 10 artists>)
```




```
[ ]: plt.hist(df['coverage'])
```

```
[ ]: (array([3462., 282., 758., 0., 297., 128., 0., 48., 21.,
          4.]),
      array([0. , 0.7, 1.4, 2.1, 2.8, 3.5, 4.2, 4.9, 5.6, 6.3, 7. ]),
      <BarContainer object of 10 artists>)
```




```
[ ]: plt.hist(df['mms'])
```

```
[ ]: (array([192., 592., 505., 581., 524., 660., 663., 501., 586., 196.]),  
      array([-3. ,  1.6,  6.2, 10.8, 15.4, 20. , 24.6, 29.2, 33.8, 38.4, 43. ]),  
      <BarContainer object of 10 artists>)
```




```
[ ]: plt.hist(df['charges'])
```

```
[ ]: (array([ 802., 1153.,  867.,  879.,  377.,  307.,  268.,  197.,  141.,
           9.]),
      array([  80.,  296.,  512.,  728.,  944., 1160., 1376., 1592., 1808.,
           2024., 2240.]),
      <BarContainer object of 10 artists>)
```




```
[ ]: plt.hist(df['sms'])
```

```
[ ]: (array([262., 487., 632., 479., 532., 600., 524., 678., 504., 302.]),  
      array([23. , 27.4, 31.8, 36.2, 40.6, 45. , 49.4, 53.8, 58.2, 62.6, 67. ]),  
      <BarContainer object of 10 artists>)
```


```
[ ]: plt.hist(df['calls'])
```

```
[ ]: (array([1683., 1376., 1039., 319., 219., 97., 132., 84., 45.,
          6.]),
      array([ 0., 100., 200., 300., 400., 500., 600., 700., 800.,
          900., 1000.]),
      <BarContainer object of 10 artists>)
```



```
[ ]: plt.hist(df['age'])
```


```
[ ]: (array([2093.,  0.,  0.,  0.,  0., 1402.,  0.,  0.,  0.,
          1499.]),
      array([1. , 1.2, 1.4, 1.6, 1.8, 2. , 2.2, 2.4, 2.6, 2.8, 3. ]),
      <BarContainer object of 10 artists>)
```



Inference: similar to count plot. Though have floating point value


```
[ ]: plt.pie(df['churn'].value_counts().values,shadow=True , autopct='%1.2f%%',
           ↳labels=df.churn.value_counts().index)
```

```
[ ]: ([<matplotlib.patches.Wedge at 0x79b033e9e390>,
      <matplotlib.patches.Wedge at 0x79b03395d8d0>],
      [Text(-1.0503509720705493, 0.32674582701306604, 'No Churn'),
       Text(1.0503510785377979, -0.3267454847652605, 'Churn')],
      [Text(-0.5729187120384813, 0.17822499655258148, '90.40%'),
       Text(0.572918770111526, -0.17822480987196024, '9.60%')])
```




```
[ ]: plt.pie(df['phone'].value_counts().values,shadow=True , autopct='%1.2f%%',
    ↪labels=df.phone.value_counts().index)
```

```
[ ]: ([<matplotlib.patches.Wedge at 0x79b033991f10>,
    <matplotlib.patches.Wedge at 0x79b0339a0590>],
    [Text(-1.0802560936468724, 0.2074771605232722, 'Andoid'),
    Text(1.080256161251455, -0.20747680853114742, 'IOS')],
    [Text(-0.5892305965346576, 0.11316936028542118, '93.96%'),
    Text(0.5892306334098845, -0.11316916828971677, '6.04%')])
```




```
[ ]: plt.pie(df['prepost'].value_counts().values,shadow=True , autopct='%1.2f%%' ,
↳labels=df.prepost.value_counts().index)
```

```
[ ]: ([<matplotlib.patches.Wedge at 0x79b0339d9450>,
      <matplotlib.patches.Wedge at 0x79b0339cdc50>],
      [Text(-0.3293846408032828, 1.0495264448325696, 'Postpaid'),
       Text(0.3293851078956752, -1.0495262982396176, 'Prepaid')],
      [Text(-0.17966434952906332, 0.5724689699086742, '59.68%'),
       Text(0.1796646043067319, -0.5724688899488822, '40.32%')])
```




```
[ ]: plt.pie(df['gender'].value_counts().values,shadow=True , autopct='%1.2f%%',  
↳labels=df.gender.value_counts().index)
```

```
[ ]: ([<matplotlib.patches.Wedge at 0x79b033a16990>,  
      <matplotlib.patches.Wedge at 0x79b033a24d10>],  
      [Text(-0.663222493530133, 0.8775738852516481, 'Male'),  
       Text(0.6632225702500791, -0.8775738272708907, 'Female')],  
      [Text(-0.36175772374370885, 0.47867666468271713, '70.60%'),  
       Text(0.3617577655909522, -0.4786766330568494, '29.40%')])
```




```
[ ]: plt.pie(df['sim'].value_counts().values,shadow=True , autopct='%1.2f%%',
↳labels=df.sim.value_counts().index)
```

```
[ ]: ([<matplotlib.patches.Wedge at 0x79b033864310>,
      <matplotlib.patches.Wedge at 0x79b033866a50>],
      [Text(-1.0413637338022248, 0.35434668605969627, 'Single Sim'),
       Text(1.0413638492629755, -0.35434634674030313, 'Dual Sim')],
      [Text(-0.5680165820739408, 0.1932800105780161, '89.56%'),
       Text(0.568016645052532, -0.19327982549471076, '10.44%')])
```




```
[ ]: sns.boxplot(x='calls', data=df)
```

```
[ ]: <Axes: xlabel='calls'>
```


```
[ ]: sns.boxplot(x='sms', data=df)
```

```
[ ]: <Axes: xlabel='sms'>
```




```
[ ]: sns.boxplot(x='mms', data=df)
```

```
[ ]: <Axes: xlabel='mms'>
```




```
[ ]: sns.boxplot(x='charges', data=df)
```

```
[ ]: <Axes: xlabel='charges'>
```



```
[ ]: sns.boxplot(x='coverage', data=df)
```

```
[ ]: <Axes: xlabel='coverage'>
```




```
[ ]: print(df['calls'].quantile(0.10))  
      print(df['calls'].quantile(0.85))
```

```
30.0  
340.0
```


```
[ ]: df['calls'] = np.where(df['calls'] > 340 , 340 , df['calls'])  
      df['calls'] = np.where(df['calls'] < 30 , 30 , df['calls'])
```

```
[ ]: sns.boxplot(data=df , x='calls')
```

```
[ ]: <Axes: xlabel='calls'>
```



```
[ ]: fig,ax = plt.subplots(figsize=(15,25))
sns.heatmap(df.corr(numeric_only = True) , annot=True , cmap='RdBu')
plt.show()
```



```
[ ]: df.drop(axis=1 , columns = ['phoneno' , 'zipcode' ] , inplace=True)
```

```
[ ]: df.churn.value_counts()
```

```
[ ]: churn
No Churn    4520
Churn       480
Name: count, dtype: int64
```

Handling Missing

```
[ ]: df['age'].fillna(df['age'].median())
```

```
[ ]: 0      1.0
      1      1.0
      2      1.0
      3      2.0
      4      2.0
      ...
      4995   3.0
      4996   1.0
      4997   3.0
      4998   2.0
      4999   1.0
      Name: age, Length: 5000, dtype: float64
```

Handling Categorical

```
[ ]: df.head()
```

```
[ ]:   age  gender  calls  sms  mms  charges  coverage  complaint  sim \
0  1.0   Male   160   25    1    490         0         4   Dual Sim
1  1.0   Male   150   45   19    340         0         3   Dual Sim
2  1.0   Male   100   39   15    110         0         1  Single Sim
3  2.0   Male   270   35    9   1000         0         1  Single Sim
4  2.0  Female   100   35    8    450         0         4  Single Sim
```

```
      phone  prepost  churn
0  Andoid  Prepaid  No Churn
1  Andoid  Prepaid  No Churn
2  Andoid  Prepaid  No Churn
3  Andoid  Prepaid  No Churn
4  Andoid  Prepaid  No Churn
```

```
[ ]: df['phone'] = df.apply(lambda x: 1 if x['phone'] == 'Android' else 0, axis=1)
```



```
[ ]: df['gender'] = df.apply(lambda x: 1 if x['gender'] == 'Female' else 0, axis=1)
```

```
[ ]: df['prepost'] = df.apply(lambda x: 1 if x['prepost'] == 'Prepaid' else 0, axis=1)
```

```
[ ]: df['sim'] = df.apply(lambda x: 1 if x['sim'] == 'Dual Sim' else 0, axis=1)
```

```
[ ]: df.head()
```

```
[ ]:
  age  gender  calls  sms  mms  charges  coverage  complaint  sim  phone  \
0  1.0      0   160   25    1    490         0           4     1     0
1  1.0      0   150   45   19    340         0           3     1     0
2  1.0      0   100   39   15    110         0           1     0     0
3  2.0      0   270   35    9   1000         0           1     0     0
4  2.0      1   100   35    8    450         0           4     0     0
```

```

  prepost  churn
0         1  No Churn
1         1  No Churn
2         1  No Churn
3         1  No Churn
4         1  No Churn
```

Normalization

```
[ ]: df['charges'] = ((df['charges'] - df['charges'].mean()) / (df['charges'].std()))
```

```
[ ]: df['calls'] = ((df['calls'] - df['calls'].mean()) / (df['calls'].std()))
```

```
[ ]: df.head()
```

```
[ ]:
  age  gender  calls  sms  mms  charges  coverage  complaint  sim  phone  \
0  1.0      0 -0.060834   25    1 -0.538175         0           4     1     0
1  1.0      0 -0.153574   45   19 -0.864023         0           3     1     0
2  1.0      0 -0.617274   39   15 -1.363657         0           1     0     0
3  2.0      0  0.959306   35    9  0.569708         0           1     0     0
4  2.0      1 -0.617274   35    8 -0.625068         0           4     0     0
```

```

  prepost  churn
0         1  No Churn
1         1  No Churn
2         1  No Churn
3         1  No Churn
4         1  No Churn
```

```
[ ]: from sklearn.model_selection import train_test_split
     from sklearn.preprocessing import StandardScaler
```

```

from sklearn.cluster import KMeans
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report

```

```

[ ]: y = df['churn']
     x = df.drop('churn' , axis=1)

```

```

[ ]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.
     ↪2,random_state=1)

```

```

[ ]: model = DecisionTreeClassifier(random_state=42)
     model.fit(x_train, y_train)

```

```

[ ]: DecisionTreeClassifier(random_state=42)

```

```

[ ]: y_pred = model.predict(x_test)

```

```

[ ]: print("\nBaseline Model (No Sampling) Performance:")
     print("Accuracy:", accuracy_score(y_test, y_pred))
     print(classification_report(y_test, y_pred))

```

Baseline Model (No Sampling) Performance:

Accuracy: 0.981

	precision	recall	f1-score	support
Churn	0.95	0.86	0.90	100
No Churn	0.98	0.99	0.99	900
accuracy			0.98	1000
macro avg	0.96	0.93	0.95	1000
weighted avg	0.98	0.98	0.98	1000

Simple Random Sampling

```

[ ]: churn_sample = df.sample(frac=0.3, random_state=42)

```

```

[ ]: x_sample = churn_sample.drop(columns=['churn'])
     y_sample = churn_sample['churn']

```

```

[ ]: x_train_srs, x_test_srs, y_train_srs, y_test_srs = train_test_split(x_sample,
     ↪y_sample, test_size=0.3, random_state=42)

```

```

[ ]: model.fit(x_train_srs, y_train_srs)
     y_pred_srs = model.predict(x_test_srs)

```

```
[ ]: print("\nSimple Random Sampling Performance:")
print("Accuracy:", accuracy_score(y_test_srs, y_pred_srs))
print(classification_report(y_test_srs, y_pred_srs))
```

Simple Random Sampling Performance:

Accuracy: 0.9711111111111111

	precision	recall	f1-score	support
Churn	0.80	0.86	0.83	37
No Churn	0.99	0.98	0.98	413
accuracy			0.97	450
macro avg	0.89	0.92	0.91	450
weighted avg	0.97	0.97	0.97	450

Stratified Sampling

```
[ ]: x_train_strat, x_test_strat, y_train_strat, y_test_strat = train_test_split(x,
↪y, train_size=0.3, stratify=y, random_state=42)
```

```
[ ]: model.fit(x_train_strat, y_train_strat)
y_pred_strat = model.predict(x_test_strat)
```

```
[ ]: print("\nStratified Sampling Performance:")
print("Accuracy:", accuracy_score(y_test_strat, y_pred_strat))
print(classification_report(y_test_strat, y_pred_strat))
```

Stratified Sampling Performance:

Accuracy: 0.97

	precision	recall	f1-score	support
Churn	0.83	0.87	0.85	336
No Churn	0.99	0.98	0.98	3164
accuracy			0.97	3500
macro avg	0.91	0.93	0.92	3500
weighted avg	0.97	0.97	0.97	3500

```
[ ]: from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
```

```
[ ]: scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
```

```
[ ]: imputer = SimpleImputer(strategy='mean')
x_scaled_imputed = imputer.fit_transform(x_scaled)

[ ]: kmeans = KMeans(n_clusters=3, random_state=42)
df['Cluster'] = kmeans.fit_predict(x_scaled_imputed)

[ ]: adaptive_sample = df.groupby('Cluster').sample(frac=0.3, random_state=42)

[ ]: x_adaptive = adaptive_sample.drop(columns=['churn', 'Cluster'])
y_adaptive = adaptive_sample['churn']

[ ]: X_train_adapt, X_test_adapt, y_train_adapt, y_test_adapt = \
    train_test_split(x_adaptive, y_adaptive, test_size=0.3, random_state=42)

[ ]: model.fit(X_train_adapt, y_train_adapt)
y_pred_adapt = model.predict(X_test_adapt)

[ ]: print("\nAdaptive Sampling Performance:")
print("Accuracy:", accuracy_score(y_test_adapt, y_pred_adapt))
print(classification_report(y_test_adapt, y_pred_adapt))
```

Adaptive Sampling Performance:

Accuracy: 0.9733333333333334

	precision	recall	f1-score	support
Churn	0.89	0.85	0.87	47
No Churn	0.98	0.99	0.99	403
accuracy			0.97	450
macro avg	0.94	0.92	0.93	450
weighted avg	0.97	0.97	0.97	450

```
[ ]: print("\nOriginal Dataset Statistics:")
print(df.describe())
print("\nSimple Random Sampling Statistics:")
print(churn_sample.describe())
print("\nStratified Sampling Statistics:")
print(pd.DataFrame(x_train_strat).describe())
print("\nAdaptive Sampling (K-Means) Statistics:")
print(adaptive_sample.describe())
```

Original Dataset Statistics:

	age	gender	calls	sms	mms	\
count	4994.000000	5000.000000	5.000000e+03	5000.000000	5000.000000	

mean	1.881057	0.294000	8.739676e-17	45.338400	20.104600
std	0.839796	0.455637	1.000000e+00	11.463166	11.467954
min	1.000000	0.000000	-1.266454e+00	23.000000	-3.000000
25%	1.000000	0.000000	-8.954939e-01	35.000000	10.000000
50%	2.000000	0.000000	-1.535738e-01	45.000000	20.000000
75%	3.000000	1.000000	7.738264e-01	55.000000	30.000000
max	3.000000	1.000000	1.608487e+00	67.000000	43.000000

	charges	coverage	complaint	sim	phone \
count	5.000000e+03	5000.000000	5000.000000	5000.000000	5000.0
mean	8.526513e-17	0.719000	2.388600	0.104400	0.0
std	1.000000e+00	1.233184	1.154061	0.305809	0.0
min	-1.428826e+00	0.000000	0.000000	0.000000	0.0
25%	-7.554070e-01	0.000000	1.000000	0.000000	0.0
50%	-2.123269e-01	0.000000	2.000000	0.000000	0.0
75%	5.262619e-01	2.000000	3.000000	0.000000	0.0
max	3.263385e+00	7.000000	4.000000	1.000000	0.0

	prepost	Cluster
count	5000.000000	5000.000000
mean	0.403200	0.999200
std	0.490589	0.946456
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	1.000000	2.000000
max	1.000000	2.000000

Simple Random Sampling Statistics:

	age	gender	calls	sms	mms \
count	1498.000000	1500.000000	1500.000000	1500.000000	1500.000000
mean	1.875167	0.284667	0.009729	45.924667	20.707333
std	0.844510	0.451406	1.004375	11.409489	11.426083
min	1.000000	0.000000	-1.266454	23.000000	-3.000000
25%	1.000000	0.000000	-0.895494	36.000000	11.000000
50%	2.000000	0.000000	-0.060834	46.000000	21.000000
75%	3.000000	1.000000	0.866566	55.000000	30.000000
max	3.000000	1.000000	1.608487	67.000000	43.000000

	charges	coverage	complaint	sim	phone	prepost
count	1500.000000	1500.000000	1500.000000	1500.000000	1500.0	1500.000000
mean	0.009626	0.702667	2.418667	0.106667	0.0	0.405333
std	1.016614	1.198562	1.151348	0.308792	0.0	0.491120
min	-1.428826	0.000000	0.000000	0.000000	0.0	0.000000
25%	-0.755407	0.000000	1.000000	0.000000	0.0	0.000000
50%	-0.212327	0.000000	2.000000	0.000000	0.0	0.000000
75%	0.591432	1.000000	4.000000	0.000000	0.0	1.000000
max	2.828921	7.000000	4.000000	1.000000	0.0	1.000000

Stratified Sampling Statistics:

	age	gender	calls	sms	mms \
count	1497.000000	1500.000000	1500.000000	1500.000000	1500.000000
mean	1.888444	0.291333	-0.032171	44.896667	19.667333
std	0.834728	0.454528	1.014331	11.089754	11.087288
min	1.000000	0.000000	-1.266454	23.000000	-3.000000
25%	1.000000	0.000000	-0.988234	35.000000	10.000000
50%	2.000000	0.000000	-0.153574	45.000000	20.000000
75%	3.000000	1.000000	0.773826	54.000000	29.000000
max	3.000000	1.000000	1.608487	66.000000	41.000000

	charges	coverage	complaint	sim	phone	prepost
count	1500.000000	1500.000000	1500.000000	1500.000000	1500.0	1500.000000
mean	-0.008476	0.718667	2.36200	0.092667	0.0	0.408667
std	1.021317	1.242798	1.14826	0.290061	0.0	0.491751
min	-1.428826	0.000000	0.00000	0.000000	0.0	0.000000
25%	-0.777130	0.000000	1.00000	0.000000	0.0	0.000000
50%	-0.255773	0.000000	2.00000	0.000000	0.0	0.000000
75%	0.553416	2.000000	3.00000	0.000000	0.0	1.000000
max	2.850645	6.000000	4.00000	1.000000	0.0	1.000000

Adaptive Sampling (K-Means) Statistics:


	age	gender	calls	sms	mms \
count	1499.000000	1500.000000	1500.000000	1500.000000	1500.000000
mean	1.895264	0.302667	0.010533	45.291333	20.045333
std	0.843002	0.459565	1.002715	11.490684	11.529089
min	1.000000	0.000000	-1.266454	23.000000	-2.000000
25%	1.000000	0.000000	-0.895494	35.000000	10.000000
50%	2.000000	0.000000	-0.060834	45.000000	20.000000
75%	3.000000	1.000000	0.866566	55.000000	30.000000
max	3.000000	1.000000	1.608487	67.000000	43.000000

	charges	coverage	complaint	sim	phone \
count	1500.000000	1500.000000	1500.000000	1500.000000	1500.0
mean	-0.000149	0.743333	2.387333	0.104667	0.0
std	1.017994	1.263271	1.163705	0.306226	0.0
min	-1.428826	0.000000	0.000000	0.000000	0.0
25%	-0.755407	0.000000	1.000000	0.000000	0.0
50%	-0.255773	0.000000	2.000000	0.000000	0.0
75%	0.461092	2.000000	3.000000	0.000000	0.0
max	2.850645	6.000000	4.000000	1.000000	0.0

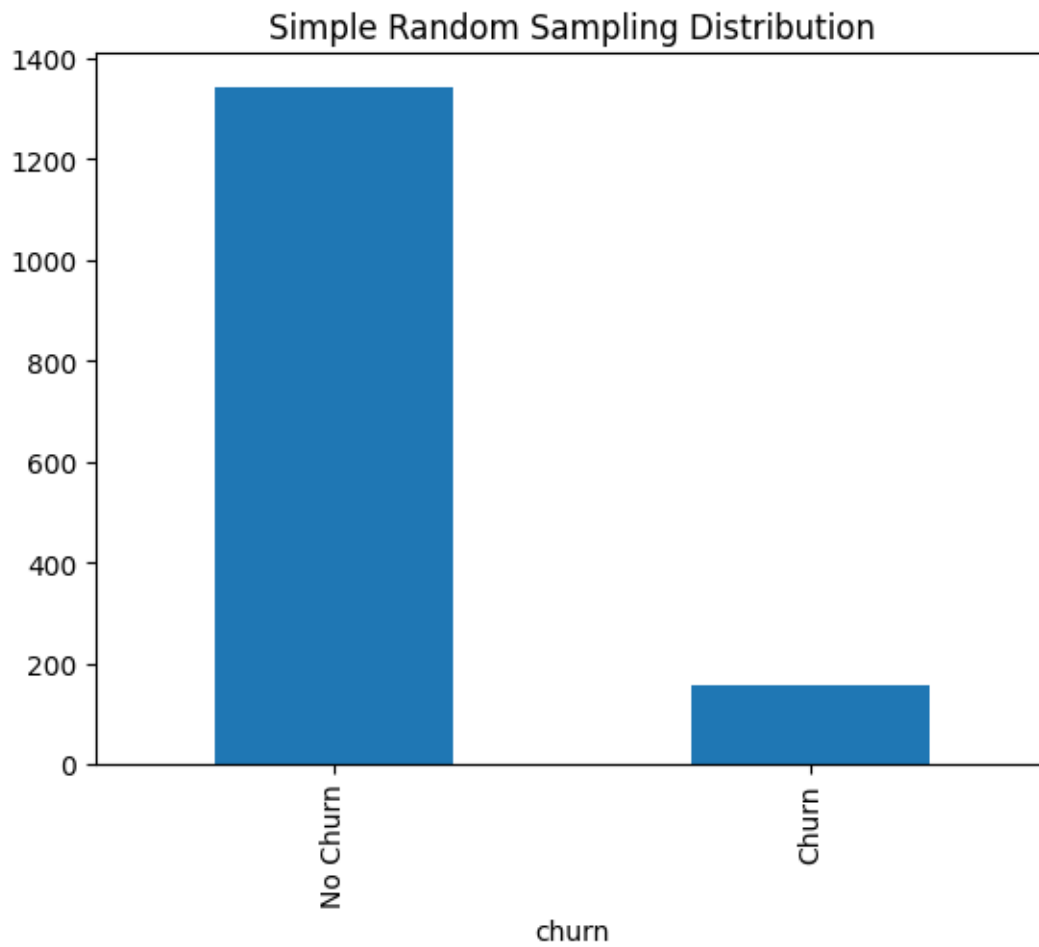
	prepost	Cluster
count	1500.000000	1500.000000
mean	0.398667	0.999333
std	0.489787	0.946536
min	0.000000	0.000000

25%	0.000000	0.000000
50%	0.000000	1.000000
75%	1.000000	2.000000
max	1.000000	2.000000

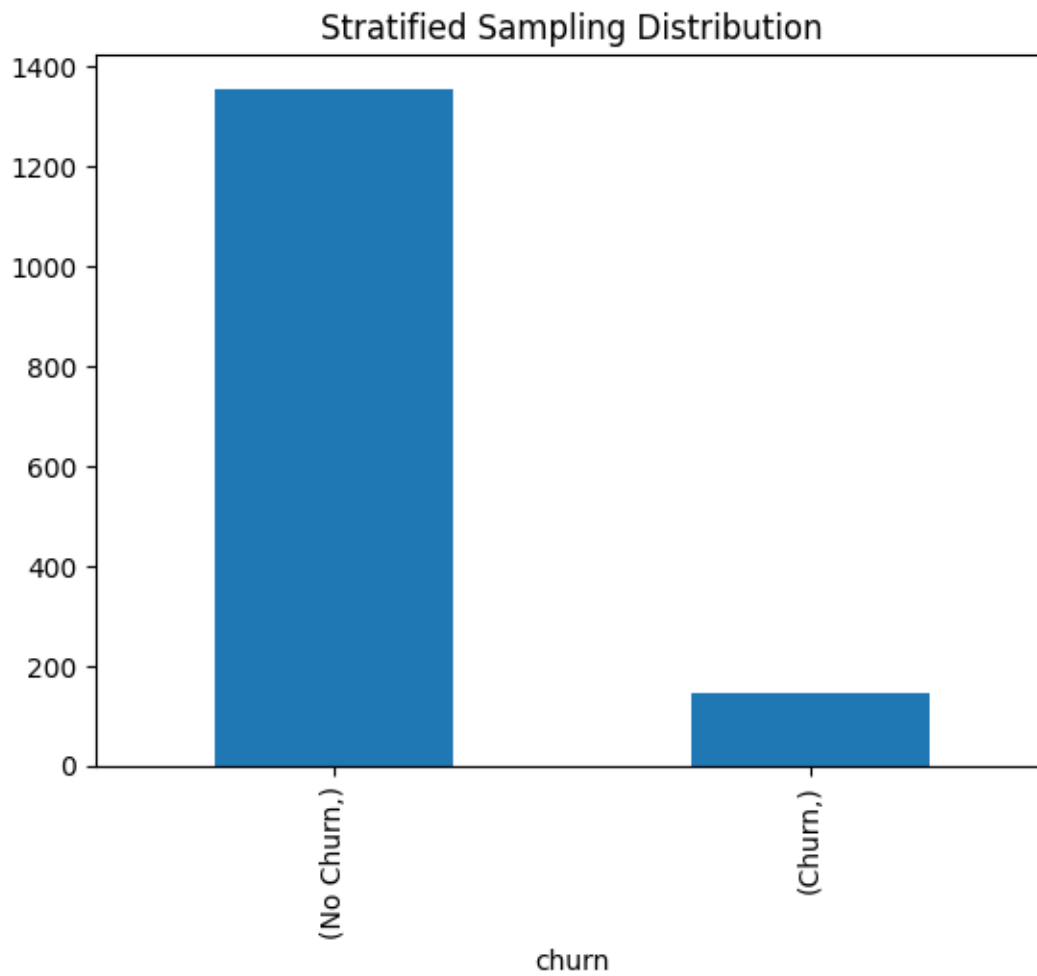
```
[ ]: import matplotlib.pyplot as plt
df['churn'].value_counts().plot(kind='bar', title='Original Dataset_
↳Distribution')
plt.show()
```



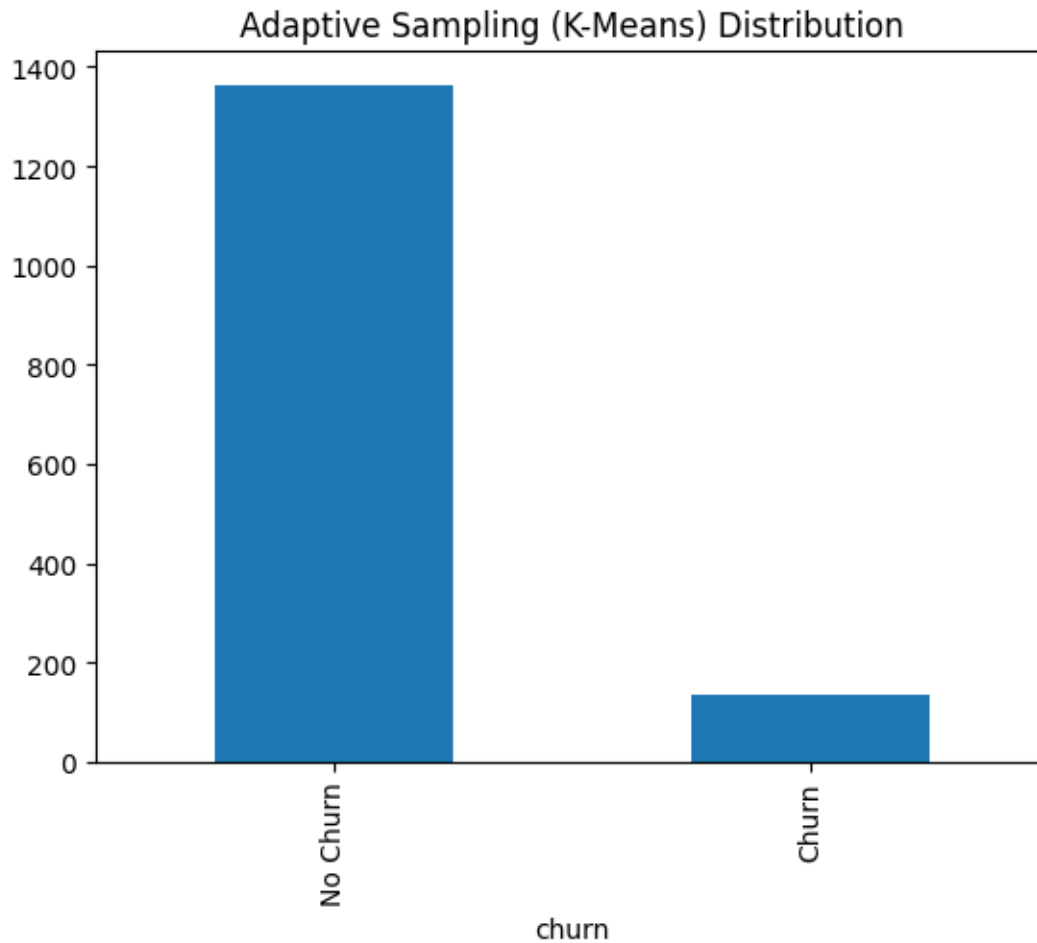
```
[ ]: churn_sample['churn'].value_counts().plot(kind='bar', title='Simple Random_
↳Sampling Distribution')
plt.show()
```



```
[ ]: pd.DataFrame({'churn': y_train_strat}).value_counts().plot(kind='bar',  
    title='Stratified Sampling Distribution')  
plt.show()
```

```
[ ]: adaptive_sample['churn'].value_counts().plot(kind='bar', title='Adaptive_␣  
      ↪Sampling (K-Means) Distribution')  
plt.show()
```



```
[ ]: import numpy as np
from sklearn.metrics import pairwise_distances

original_variance = np.var(x, axis=0)
random_variance = np.var(x_sample, axis=0)
stratified_variance = np.var(x_train_strat, axis=0)
adaptive_variance = np.var(x_adaptive, axis=0)
print("\nFeature Variance (Original Dataset):", original_variance)
print("Feature Variance (Random Sampling):", random_variance)
print("Feature Variance (Stratified Sampling):", stratified_variance)
print("Feature Variance (Adaptive Sampling):", adaptive_variance)
adaptive_spread = np.mean(pairwise_distances(x_scaled_imputed)[kmeans.labels_␣
↪ == kmeans.labels_[ :, None]])
print("\nAdaptive Sampling (K-Means) Cluster Spread:", adaptive_spread)
```

Feature Variance (Original Dataset): age

0.705116

gender	0.207564	
calls	0.999800	
sms	131.377885	
mms	131.487659	
charges	0.999800	
coverage	1.520439	
complaint	1.331590	
sim	0.093501	
phone	0.000000	
prepost	0.240630	
dtype: float64		
Feature Variance (Random Sampling): age		0.712721
gender	0.203632	
calls	1.008096	
sms	130.089658	
mms	130.468346	
charges	1.032814	
coverage	1.435593	
complaint	1.324718	
sim	0.095289	
phone	0.000000	
prepost	0.241038	
dtype: float64		
Feature Variance (Stratified Sampling): age		0.696306
gender	0.206458	
calls	1.028182	
sms	122.900656	
mms	122.846000	
charges	1.042394	
coverage	1.543518	
complaint	1.317623	
sim	0.084080	
phone	0.000000	
prepost	0.241658	
dtype: float64		
Feature Variance (Adaptive Sampling): age		0.710178
gender	0.211060	
calls	1.004767	
sms	131.947792	
mms	132.831278	
charges	1.035622	
coverage	1.594789	
complaint	1.353306	
sim	0.093712	
phone	0.000000	
prepost	0.239732	
dtype: float64		

Adaptive Sampling (K-Means) Cluster Spread: 3.757266708546486